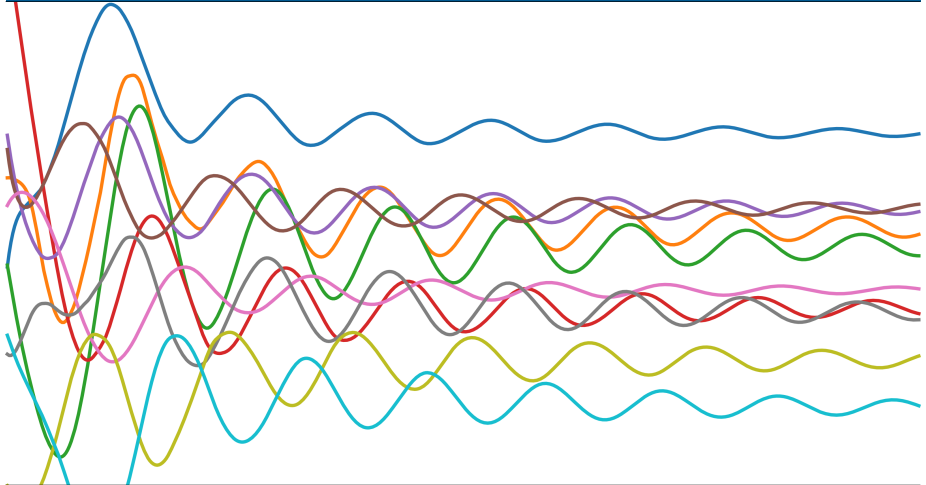


Variational Autoencoders for Koopman Dynamical Systems



TECHNISCHE
UNIVERSITÄT
DARMSTADT

B.Sc. Intermediate Presentation



Motivation

- Control theory for linear systems is highly evolved.
- But nonlinear systems are hard...
- We need ways to linearize a nonlinear system!

Motivation

- Control theory for linear systems is highly evolved.
- But nonlinear systems are hard...
- We need ways to linearize a nonlinear system!
- “Classical” Linearization:
 - ▣ Use small angle approximation (e.g. for pendulum)?
 - ▣ Diverge fast with higher displacements...
 - ▣ Only linearize locally, not globally.

Motivation

- Control theory for linear systems is highly evolved.
- But nonlinear systems are hard...
- We need ways to linearize a nonlinear system!
- “Classical” Linearization:
 - ▢ Use small angle approximation (e.g. for pendulum)?
 - ▢ Diverge fast with higher displacements...
 - ▢ Only linearize locally, not globally.
- Koopman theory offers a way to do that!

The Koopman Operator

For a *nonlinear* dynamical system

$$\mathbf{s}_{t+1} = \mathbf{F}(\mathbf{s}_t)$$

with *nonlinear* measurements (and *embedding*)


$$\mathbf{y}_t = \mathbf{g}(\mathbf{s}_t)$$

the Koopman operator advances these measurements forward in time *linearly*:

$$\mathbf{g}(\mathbf{s}_{t+1}) = \mathcal{K}\mathbf{g}(\mathbf{s}_t) \quad \Longleftrightarrow \quad \mathbf{y}_{t+1} = \mathcal{K}\mathbf{y}_t$$


This is possible for every nonlinear dynamical system. And globalizes linearly! But the embedding \mathbf{g} is typically infinite-dimensional...

Koopman Dynamical System (“Deterministic”)



Adopted from Brunton et al. “Koopman Invariant Subspaces and Finite Linear Representations of Nonlinear Dynamical Systems for Control”.

Linear Gaussian Dynamical System



$$\mathbf{s}_{t+1} = \mathbf{A}\mathbf{s}_t + \mathbf{v}, \quad \mathbf{v} \sim \mathcal{N}(\mathbf{0}, \mathbf{Q})$$

$$\mathbf{y}_t = \mathbf{g}_\theta(\mathbf{s}_t) + \mathbf{w}, \quad \mathbf{w} \sim \mathcal{N}(\mathbf{0}, \mathbf{R})$$



$$\mathbf{s}_{t+1} \sim \mathcal{N}(\mathbf{A}\mathbf{s}_t, \mathbf{Q})$$

$$\mathbf{y}_t \sim \mathcal{N}(\mathbf{g}_\theta(\mathbf{s}_t), \mathbf{R})$$

How to learn?

Goal: Estimate all of the following:

- Latent dynamics matrix A .
- Measurement function $g_{\theta}(\cdot)$ (i.e. the parameters θ).
Variational auto-encoder without an amortization network.
- Noise covariances Q, R .

We employ an EM-algorithm to do that.

How to learn?

Goal: Estimate all of the following:

- Latent dynamics matrix A .
- Measurement function $g_{\theta}(\cdot)$ (i.e. the parameters θ).
Variational auto-encoder without an amortization network.
- Noise covariances Q, R .

We employ an EM-algorithm to do that.

Core Problem

Some expectations cannot be evaluated in closed form for a nonlinear $g_{\theta}(\cdot)$.

How to learn?

Goal: Estimate all of the following:

- Latent dynamics matrix A .
- Measurement function $g_{\theta}(\cdot)$ (i.e. the parameters θ).
Variational auto-encoder without an amortization network.
- Noise covariances Q, R .

We employ an EM-algorithm to do that.

Core Problem

Some expectations cannot be evaluated in closed form for a nonlinear $g_{\theta}(\cdot)$.

- Use cubature rules for evaluating the intractable expectations.
- No closed form solution for maximizing w.r.t. function parameters θ .

How to learn?

Goal: Estimate all of the following:

- Latent dynamics matrix A .
- Measurement function $g_{\theta}(\cdot)$ (i.e. the parameters θ).
Variational auto-encoder without an amortization network.
- Noise covariances Q, R .

We employ an EM-algorithm to do that.

Core Problem

Some expectations cannot be evaluated in closed form for a nonlinear $g_{\theta}(\cdot)$.


- Use cubature rules for evaluating the intractable expectations.
- No closed form solution for maximizing w.r.t. function parameters θ .
→ Use backpropagation and gradient descent!

Results: Damped Inverted Pendulum Dynamical System

$$\ddot{\varphi} = \frac{g}{L} \sin(\varphi) - d\dot{\varphi} \quad \longrightarrow \quad \ddot{\varphi} = \sin(\varphi) - 0.1\dot{\varphi}$$

Observations:

- Displacement φ
- Velocity $\dot{\varphi}$

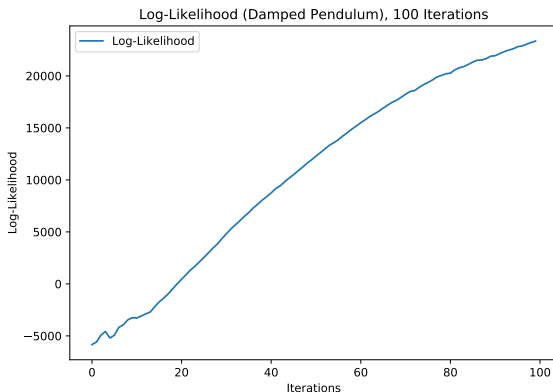



Results: Damped Inverted Pendulum Dynamical System

$$\ddot{\varphi} = \frac{g}{L} \sin(\varphi) - d\dot{\varphi} \quad \longrightarrow \quad \ddot{\varphi} = \sin(\varphi) - 0.1\dot{\varphi}$$

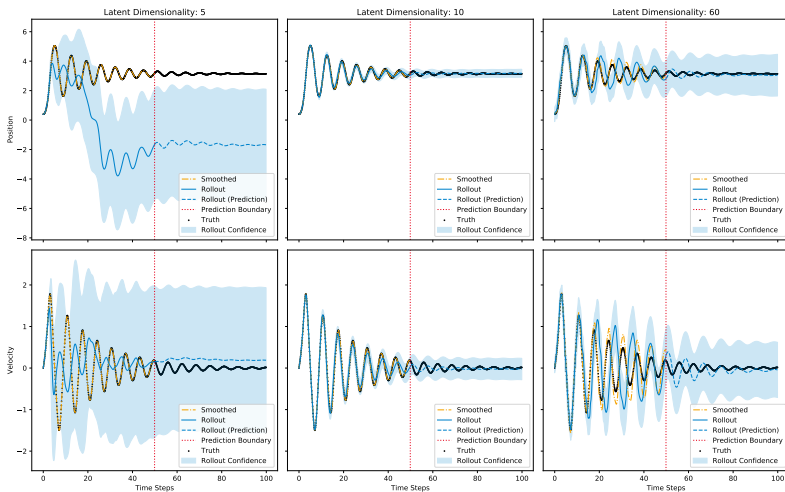
Observations:

- Displacement φ
- Velocity $\dot{\varphi}$



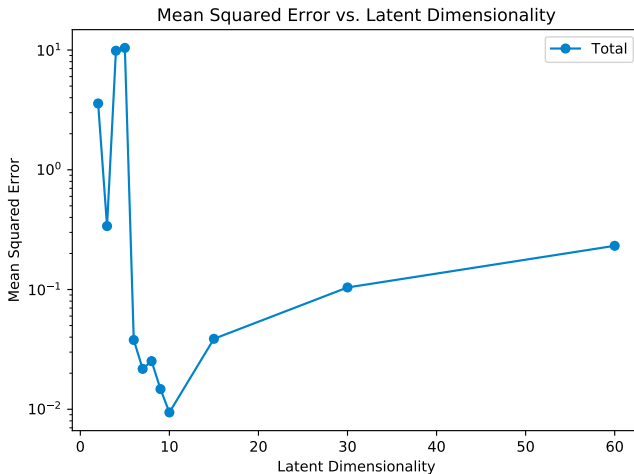
Results: Damped Inverted Pendulum

Rollout/Prediction in Observation Space, 10 La- tents



Results: Damped Inverted Pendulum

Different Sizes of Latent Dimension



Related Work

- Lusch et al. “Deep Learning for Universal Linear Embeddings of Nonlinear Dynamics”:
 - ▣ Use an autoencoder approach with an encoder network.
 - ▣ No probabilistic interpretation.

Related Work

- Lusch et al. “Deep Learning for Universal Linear Embeddings of Nonlinear Dynamics”:
 - ▣ Use an autoencoder approach with an encoder network.
 - ▣ No probabilistic interpretation.
- Morton et al. “Deep Variational Koopman Models: Inferring Koopman Observations for Uncertainty-Aware Dynamics Modeling and Control”:
 - ▣ Utilize six different neural networks to perform variational inference.
 - ▣ We use only one!
 - ▣ And we expect to do at least as good.

Related Work

- Lusch et al. “Deep Learning for Universal Linear Embeddings of Nonlinear Dynamics”:
 - ▣ Use an autoencoder approach with an encoder network.
 - ▣ No probabilistic interpretation.
- Morton et al. “Deep Variational Koopman Models: Inferring Koopman Observations for Uncertainty-Aware Dynamics Modeling and Control”:
 - ▣ Utilize six different neural networks to perform variational inference.
 - ▣ We use only one!
 - ▣ And we expect to do at least as good.
- Becker et al. “Recurrent Kalman Networks: Factorized Inference in High-Dimensional Deep Feature Spaces”

Related Work

- Lusch et al. “Deep Learning for Universal Linear Embeddings of Nonlinear Dynamics”:
 - ▢ Use an autoencoder approach with an encoder network.
 - ▢ No probabilistic interpretation.
- Morton et al. “Deep Variational Koopman Models: Inferring Koopman Observations for Uncertainty-Aware Dynamics Modeling and Control”:
 - ▢ Utilize six different neural networks to perform variational inference.
 - ▢ We use only one!
 - ▢ And we expect to do at least as good.
- Becker et al. “Recurrent Kalman Networks: Factorized Inference in High-Dimensional Deep Feature Spaces”
- Zhang, Vikram et al. “SOLAR: Deep Structured Representations for Model-Based Reinforcement Learning”

Hypothesis and Outlook

- Tackle the ELBO using approximate EM rather than stochastic variational inference.
- Assume we can do at least as good as Morton et al. with simpler network architectures.
- Perform control tasks on pendulum, cartpole, acrobot.

Backup Slides

Outline

Backup Slides

- Expectation Maximization

- The Expected Log-Likelihood

- Cubature Rules

Expectation Maximization

- E-Step: Calculate the expected latents and correlations using filtering/smoothing.
- M-Step: Maximize the expected log-likelihood $\mathbb{E}[\ln p(s_{1:T}, y_{1:T}) | y_{1:T}]$.

The Expected Log-Likelihood

Markov property yields complete log-likelihood:

$$\ln p(\mathbf{s}_{1:T}, \mathbf{y}_{1:T}) = \ln p(\mathbf{s}_1) + \sum_{t=2}^T \ln p(\mathbf{s}_{t+1} | \mathbf{s}_t) + \sum_{t=1}^T \ln p(\mathbf{y}_t | \mathbf{s}_t)$$

Expectation $\mathbb{E}[\ln p(\mathbf{s}_{1:T}, \mathbf{y}_{1:T}) | \mathbf{y}_{1:T}]$ is based on five other expectations:

$$\hat{\mathbf{s}}_t := \mathbb{E}[\mathbf{s}_t | \mathbf{y}_{1:T}] \quad \mathbf{P}_t := \mathbb{E}[\mathbf{s}_t \mathbf{s}_t^T | \mathbf{y}_{1:T}] \quad \mathbf{P}_{t,t-1} := \mathbb{E}[\mathbf{s}_t \mathbf{s}_{t-1}^T | \mathbf{y}_{1:T}]$$

$$\hat{g}_t := \mathbb{E}[g(\mathbf{s}_t) | \mathbf{y}_{1:T}] \quad \mathbf{G}_t := \mathbb{E}[g(\mathbf{s}_t) g^T(\mathbf{s}_t) | \mathbf{y}_{1:T}]$$

Quadrature for High Dimensions: Cubature Rules

The Spherical-Radial Cubature Rule

Approximation of an arbitrary Gaussian expectation:

$$\begin{aligned}\mathbb{E}_{\mathbf{x} \sim \mathcal{N}(\boldsymbol{\mu}, \boldsymbol{\Sigma})} [\mathbf{f}(\mathbf{x})] &= \int_{\mathbb{R}^n} \mathbf{f}(\mathbf{x}) \mathcal{N}(\mathbf{x} | \boldsymbol{\mu}, \boldsymbol{\Sigma}) \, d\mathbf{x} \\ &\approx \frac{1}{2n} \sum_{i=1}^{2n} \mathbf{f}(\sqrt{\boldsymbol{\Sigma}} \boldsymbol{\xi}_i + \boldsymbol{\mu}), \quad \boldsymbol{\xi}_i = \sqrt{n} [\mathbf{1}]_i\end{aligned}$$

This finite sum can be evaluated!